

Sentiment Analysis Dashboard Cost Analysis Estimate Report

Service Overview

Sentiment Analysis Dashboard is a fully managed, serverless service that allows you to This project uses multiple AWS services.. This service follows a pay-as-you-go pricing model, making it cost-effective for various workloads.

Pricing Model

This cost analysis estimate is based on the following pricing model: -
ON DEMAND pricing (pay-as-you-go) unless otherwise specified -
Standard service configurations without reserved capacity or savings plans - No caching or optimization techniques applied

Assumptions

- Standard ON DEMAND pricing model
- Using Claude 3.5 Haiku for sentiment analysis (cost-effective option)
- Average feedback text length of 200 characters (~50 tokens)
- Average response length of 100 characters (~25 tokens)
- Lambda functions with 512 MB memory allocation
- DynamoDB on-demand billing mode
- API Gateway REST API pricing
- US East (N. Virginia) region pricing

Limitations and Exclusions

- Data transfer costs between regions
- CloudWatch logging and monitoring costs
- Development and maintenance costs
- React frontend hosting costs (local development)
- SSL certificate costs
- Custom domain costs

Cost Breakdown

Unit Pricing Details

Service	Resource Type	Unit	Price	Free Tier
Amazon Bedrock (Claude 3.5 Haiku)	Input Tokens	1,000 tokens	\$0.00025	No free tier for Bedrock foundation models

Amazon Bedrock (Claude 3.5 Haiku)	Output Tokens	1,000 tokens	\$0.00125	No free tier for Bedrock foundation models
AWS Lambda	Requests	1,000,000 requests	\$0.20	First 12 months: 1M requests/month and 400,000 GB-seconds/month free
AWS Lambda	Compute	GB-second	\$0.0000166667	First 12 months: 1M requests/month and 400,000 GB-seconds/month free
Amazon API Gateway	Requests	million requests (first 333M)	\$3.50	No free tier for API Gateway
Amazon DynamoDB	Write Requests	million write request write requests	\$0.625	First 12 months: 25 GB storage, 25 WCU, 25 RCU free
Amazon DynamoDB	Read Requests	million read request read requests	\$0.125	First 12 months: 25 GB storage, 25 WCU, 25 RCU free
Amazon DynamoDB	Storage	GB-month (after 25GB free tier)	\$0.25	First 12 months: 25 GB storage, 25 WCU, 25 RCU free

Cost Calculation

Service	Usage	Calculation	Monthly Cost
Amazon Bedrock (Claude 3.5 Haiku)	Processing sentiment analysis requests with Claude 3.5 Haiku (Input Tokens: 500,000 tokens) + \$0.00125/1K × 250K output tokens	\$0.00025/1K × 500K input tokens + \$0.00125/1K × 250K output tokens = \$0.125 + \$0.3125 = \$0.4375 per month for 10K	\$0.19

	avg), requests	
	Output	
	Tokens:	
	250,000	
	tokens	
	(10,000	
	requests ×	
	25 tokens	
	avg))	
	2 Lambda	
	functions:	
	sentiment	
	analysis	
	processor	
	and data	
	retrieval	
	(Requests:	
	20,000	$\$0.20/1M \times 0.02M$
	requests	
AWS	(10K	$\$0.0000166667 \times$
Lambda	analyze +	20,000 GB-seconds
	10K	$= \$0.004 + \0.333
	retrieve),	$= \$0.337$ per
	Compute:	month
	20,000	
	requests ×	
	2s avg ×	
	0.5GB =	
	20,000	
	GB-	
	seconds)	
	REST API	
	with 2	
	endpoints	
	for	
	feedback	
	submission	$\$3.50/1M \times 0.02M$
	and	$\$0.07$
	retrieval	$\$0.07$ per month
Amazon API	(Requests:	
Gateway	20,000 API	
	requests	
	per month)	
	Single	
	table	
	storing	
	feedback	
	analysis	
	results	
	with on-	
	demand	
	billing	$\$0.625/1M \times$
	(Write	$0.01M$ writes +
	Requests:	$\$0.125/1M \times$
	10,000	$0.01M$ reads +
	write	$\$0.25 \times 0GB$ (free tier) = $\$0.00625 +$
Amazon	requests	$\$0.00125 + \$0 =$
DynamoDB	per month,	$\$0.0075$ per month
	Read	

Requests: 10,000 read requests per month, Storage: 1 GB per month (within free tier))	Total	All services	Sum of all calculations	\$2.20/month
--	--------------	-------------------------	------------------------------------	---------------------

Free Tier

Free tier information by service:

- **Amazon Bedrock (Claude 3.5 Haiku)**: No free tier for Bedrock foundation models
- **AWS Lambda**: First 12 months: 1M requests/month and 400,000 GB-seconds/month free
- **Amazon API Gateway**: No free tier for API Gateway
- **Amazon DynamoDB**: First 12 months: 25 GB storage, 25 WCU, 25 RCU free

Cost Scaling with Usage

The following table illustrates how cost estimates scale with different usage levels:

Service	Low Usage	Medium Usage	High Usage
Amazon Bedrock (Claude 3.5 Haiku)	\$0/month	\$0/month	\$0/month
AWS Lambda	\$0/month	\$0/month	\$0/month
Amazon API Gateway	\$0/month	\$0/month	\$0/month
Amazon DynamoDB	\$0/month	\$1/month	\$3/month

Key Cost Factors

- **Amazon Bedrock (Claude 3.5 Haiku)**: Processing sentiment analysis requests with Claude 3.5 Haiku
- **AWS Lambda**: 2 Lambda functions: sentiment analysis processor and data retrieval
- **Amazon API Gateway**: REST API with 2 endpoints for feedback submission and retrieval
- **Amazon DynamoDB**: Single table storing feedback analysis results with on-demand billing

Projected Costs Over Time

The following projections show estimated monthly costs over a 12-month period based on different growth patterns:

Base monthly cost calculation:

Service	Monthly Cost
----------------	---------------------

Amazon Bedrock (Claude 3.5 Haiku)	\$0.19
AWS Lambda	\$0.38
Amazon API Gateway	\$0.07
Amazon DynamoDB	\$1.56
Total Monthly Cost	\$2

Growth Pattern	Month 1	Month 3	Month 6	Month 12
Steady	\$2/mo	\$2/mo	\$2/mo	\$2/mo
Moderate	\$2/mo	\$2/mo	\$2/mo	\$3/mo
Rapid	\$2/mo	\$2/mo	\$3/mo	\$6/mo

- Steady: No monthly growth (1.0x)
- Moderate: 5% monthly growth (1.05x)
- Rapid: 10% monthly growth (1.1x)

Detailed Cost Analysis

Pricing Model

ON DEMAND

Exclusions

- Data transfer costs between regions
- CloudWatch logging and monitoring costs
- Development and maintenance costs
- React frontend hosting costs (local development)
- SSL certificate costs
- Custom domain costs

Recommendations

Immediate Actions

- Use Claude 3.5 Haiku for cost-effective sentiment analysis
- Implement response caching for common sentiment patterns
- Configure DynamoDB on-demand billing for variable workloads
- Use ARM-based Lambda functions for 20% cost savings #### Best Practices
- Monitor token usage and optimize prompt engineering
- Implement batch processing for high-volume scenarios
- Consider provisioned throughput for predictable workloads
- Set up CloudWatch alarms for cost monitoring

Cost Optimization Recommendations

Immediate Actions

- Use Claude 3.5 Haiku for cost-effective sentiment analysis
- Implement response caching for common sentiment patterns
- Configure DynamoDB on-demand billing for variable workloads

Best Practices

- Monitor token usage and optimize prompt engineering
- Implement batch processing for high-volume scenarios
- Consider provisioned throughput for predictable workloads

Conclusion

By following the recommendations in this report, you can optimize your Sentiment Analysis Dashboard costs while maintaining performance and reliability. Regular monitoring and adjustment of your usage patterns will help ensure cost efficiency as your workload evolves.